

# Scheduling Real-Time Tasks in Multiprocessor and Distributed Systems

## Real-Time Systems Design (CS 6414)

Sumanta Pyne

Assistant Professor  
Computer Science and Engineering Department  
National Institute of Technology Rourkela  
[pynes@nitrkl.ac.in](mailto:pynes@nitrkl.ac.in)

February 26, 2021

# Overview

- Multiprocessor & distributed systems widely for real-time applications
- Recent price drop of these systems
- Dual processor architectures available at 50-60K INR even cheaper
- Distributed platform - networked PCs are common
- Better for faster response times and fault tolerance
- Distributed processing suitable for geographically distributed places
- Automated petroleum refinery - plant spread over geographic area
- Scheduling tasks more difficult than on a uniprocessor
- Uniprocessor - optimal schedule independent tasks in polynomial time
- Multiprocessor & Distributed - finding optimal schedule is NP-Hard

# Multiprocessors & Distributed Systems

- Multiprocessors tightly coupled - shared physical memory
- Distributed Systems loosely coupled - no shared memory
- Interprocess communication (IPC) expensive in tightly-coupled syst.
- IPC time ignored compared to task execution times
- Inter-task communication - read/write on shared memory
- Distributed system IPC times comparable to task execution times
- Multiprocessor uses centralized scheduler/dispatcher
- Centralized scheduler maintain task state in centralized data structure
- High communicational overheads - update data structure state change

# Scheduling on Multiprocessors & Distributed Systems

- Two subproblems
  - task allocation to processors
  - scheduling tasks on individual processors
- Task assignment problem
  - how to partition a set of tasks
  - then how to assign these to processors
  - Static - task allocation to nodes is permanent
  - Dynamic - tasks are assigned to nodes as they arise
    - different task instance allocated to different nodes
- Scheduling tasks on individual processors
  - Scheduling problem reduced to uniprocessor scheduling

# Task allocation in Multiprocessors and Distributed Systems

- An NP-Hard Problem
- Exponential time required to determine an optimal schedule
- Usage of heuristic algorithms
- Task allocation algorithms - static and dynamic
- Static algorithms
  - all tasks partitioned into subsystems
  - each subsystem is assigned a separate processor
- Dynamic algorithms
  - tasks ready for execution placed in a common priority queue
  - dispatched to processors for execution as processors become available
  - different periodic task instances execute on different processors
- Most hard real-time systems till date are static in nature
- Dynamic real-time system make more efficient resource utilization

# Multiprocessor Static Task Allocation

- Not applicable for distributed systems
- IPC time is same as memory access time due to shared memory
- Utilization Balancing algorithm
  - Maintains tasks in a queue in increasing order of utilizations
  - Removes tasks one by one from head
  - Allocates them to least utilized processor each time
  - To balance utilization of different processors
  - Perfectly balanced systems utilization  $u_i$  per processor equals overall
  - For task set  $ST_i$  assigned to processor  $P_i$ ,  $u_i = \sum_{t_j \in ST_i} u_{t_j}$
  - $u_{t_j}$  utilization due to  $T_j$
  - PR set of all processors, then total utilization is  $\sum u_i$ ,  $P_i \in PR$
  - Difficult to achieve perfect balancing of utilizations, i.e.  $u_i = u'$  for  $P_i$
  - Simple heuristic gives suboptimal results
  - Objective of good utilization balancing - Minimize  $\sum |u' - u_i|$ ,  $1 \leq i \leq n$
  - $n$  #processors,  $u'$  average utilization of processors,  $u_i$  utilization of  $P_i$
  - Suitable when #processors is fixed
  - Used when tasks at individual processors are scheduling using EDF

# Next-Fit Algorithm for RMA

- A task set is partitioned, each scheduled on a uniprocessor using RMA
- Attempts to use as few processors as possible
- Classifies different tasks into few classes based on task utilization
- One or more processors assigned to each class of tasks
- Task with similar utilization values scheduled on same processor
- Policy for utilization based task classification
  - If tasks divided to  $m$  classes,  $T_i \in \text{class } j, 0 \leq j \leq m$
- Partitioning tasks of a system into four classes
  - Class 1:  $(2^{\frac{1}{2}} - 1) < C_1 \leq (2^{\frac{1}{1}} - 1)$
  - Class 2:  $(2^{\frac{1}{3}} - 1) < C_2 \leq (2^{\frac{1}{2}} - 1)$
  - Class 3:  $(2^{\frac{1}{4}} - 1) < C_3 \leq (2^{\frac{1}{3}} - 1)$
  - Class 4:  $0 < C_4 \leq (2^{\frac{1}{4}} - 1)$
- Utilization grid for different classes
  - class 1:(0.41,1), class 2:(0.26,0.41), class 3:(0.19,0.26), class 4:(0,0.19)
- Higher task utilization values are coarser compared to lower
- Grid size: class 1 tasks is  $1 - 0.41 = 0.59$ , class 3 tasks is 0.7
- Simulation - next-fit algorithm atmost  $2.34 \times$  optimum #processors

## Example 1

The following table shows the execution times (in msec) and periods in (msec) of a set of 10 periodic real-time tasks. Assume that the tasks need to run on a multiprocessor with four processors. Allocate the tasks to processors using the next fit algorithm. Assume that the individual processors are to be scheduled using RMA algorithm.

Task Set

Task	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$	$T_6$	$T_7$	$T_8$	$T_9$	$T_{10}$
$e_i$	5	7	3	1	10	16	1	3	9	17
$p_i$	10	21	22	24	30	40	50	55	70	100

Solution

Task	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$	$T_6$	$T_7$	$T_8$	$T_9$	$T_{10}$
$e_i$	5	7	3	1	10	16	1	3	9	17
$p_i$	10	21	22	24	30	40	50	55	70	100
$u_i$	0.5	0.33	0.14	0.04	0.33	0.4	0.02	0.05	0.13	0.17
Class	1	2	4	4	2	2	4	4	4	4



# Bin Packing for EDF

- Allocates tasks to processors
- Tasks on individual processors scheduled using EDF
- Tasks assigned to  $P_i$  such that  $u_i \leq 1$
- Formulate task allocation as bin packing problem
  - Given  $n$  periodic real-time tasks
  - Individual processors scheduled using EDF
  - #bins necessary =  $\sum u_i, 1 \leq i \leq n$
  - Bin packing is NP-complete
- First-fit random bin packing algorithm
  - Tasks selected randomly
  - Assigned processors arbitrarily as long as  $u_i \leq 1$
  - Atmost  $1.7 \times$  optimum #processors required
- First-fit decreasing bin packing algorithm
  - Tasks sorted in non-increasing order of CPU utilization in ordered list
  - Task selected one by one from ordered list
  - Assigned to bin(processor  $P_i$ ) to which it fits first (i.e.  $u_i \leq 1$ )
  - Simulations show #processors required is  $1.22 \times$  optimal #processors

# Dynamic Allocation of Tasks

- Applications where tasks arrive sporadically at different nodes
- Dynamic algorithms needed to handle such tasks
- Assume any task can be executed on any processor
- Dynamic solutions naturally distributed
- No central allocation policy running on some processor
- No preallocation of tasks to processors
- Assign tasks when they arise
- Task allocation made on instantaneous load positions of nodes
- Achievable schedulable utilization better than static approaches
- High run time overhead - allocator component running at each node
- Static allocation - task assigned permanently during initialization
- No runtime overhead for static
- Dynamic ineffective if task bound to single or subset of processors

# Focussed Addressing and Bidding

- Every processor maintain status table and system load table
- Status table contains
  - information of tasks committed to run
  - execute time of tasks
  - periods of tasks
- System load table contains latest load information other processors
- Determine surplus computing capacity available
- Time axis divided into windows - intervals of fixed duration
- At end of each window each processor broadcasts to all other
- Fraction of computing power currently free for next window
- Fraction of next window with no committed tasks
- On receiving this every processor updates system load table
- When task arises, node checks whether process it locally
- If possible, it updates its status table
- Otherwise, looks out for a process can offload task

# Focussed Addressing and Bidding

- Processor finds suitable processor consulting system load table
- Determines least loaded processors to accommodate the task
- It then sends request for bids (RFBs) to these processors
- While looking for processor an overloaded processor
- Checks surplus information and selects a *focussed* processor
- information in system load table may be out of date
- Obsolete information problem
- Solution - send RFBs only if it determines that task complete in time
- High communication overhead to maintain system load tables
- Larger window size lower overhead

# Buddy Algorithm

- Aims to overcome high communication overhead
- Processor states - underloaded and overloaded
- $P_i$  underloaded if  $u_i < Th$
- $P_i$  overloaded if  $u_i \geq Th$
- Processor broadcasts - change in state
- Broadcasts to limited subset *buddy* set

# Fault-Tolerant Scheduling of Tasks

- Task scheduling can be used for fault-tolerance in real-time systems
- An effective technique
- Requires very little redundant hardware resource
- Achieved by scheduling additional ghost copies with primary task copy
- Ghost copies (clones)
  - may not be identical to primary copy
  - copies stripped down versions
  - can be executed in shorter durations than primary
- Ghost copies of different tasks can be overloaded on same slot
- Success execution of primary, deallocates corresponding backup

# Clocks in Distributed Real-Time Systems

- Clocks used for determining timeouts and time stamping
- Timeouts
  - determine failure of tasks due to deadline miss
  - indicate transmission faults or delays, or non-existent receivers
- Time stamping
  - used in message communication among tasks
  - sender includes current time along a message
  - gives age of message
  - used for ordering purposes
  - relies on good real-time clock services
- Distributed system has one clock at each node
- Different clocks diverge
- Two clocks of run exactly at same speed is impossible

# Clocks in Distributed Real-Time Systems

- Lack of synchrony among clocks expressed as *clock slew*
- *clock slew* determines attendant drift of clocks with time
- Lack of synchrony and drift - meaningless time stamping and timeout
- Synchronized clocks for meaningful time stamping and timeout
- Synchronization very important in distributed real-time systems




# Clock Synchronization

- Makes all clocks in network agree on their time values
- Different clocks agree on some time, differ from world time standard
- World time standard called Universal coordinated time (UTC)
- UTC based on international atomic time (TAI)
- TAI maintained at Paris by averaging atomic clocks around the world
- UTC signals used through GPS receivers and specialized radio stations
- Internal clock synchronization - all clocks synchronized wrt one clock
- External clock sync - a set of clocks synchronized with external clock
- Two main approaches for internal clock synchronization
  - centralized clock synchronization
  - distributed clock synchronization

# Centralized Clock Synchronization

- One of the clocks is designated as master clock or *time server*
- Other clocks of the system are slaves
- Slaves are kept in synchronization with master clock
- Slave clocks  $C_1, \dots, C_n$  are to be synchronized with master clock



- Server broadcasts its time to all clocks after every  $\Delta T$  time interval
- On receiving a broadcast, slaves set clock as per time at master clock
- $\Delta T$  should be chosen carefully
- If  $\Delta T$  too small
  - frequent broadcasts from master
  - good synchronization between slaves and master
  - high communication overhead
- If  $\Delta T$  too large, clocks may drift too much apart

# Centralized Clock Synchronization

- Assume maximum rate of drift between to individual clocks is  $\rho$
- Clock manufactures provides  $\rho$  as a specification parameter of a clock
- $\rho$  is unit less, it measures drift (time) per unit time
- Suppose clocks are resynchronized after every  $\Delta T$  interval
- Ignoring communication time required for broadcast
- Once broadcast received, clocks set to received time instantly
- Drift of any clock from master clock is bounded by  $\rho\Delta T$
- Maximum drift between any two clocks is limited to  $2\rho\Delta T$
- In reality it takes a finite amount of time to set a clock
- Suitable communication time and clock setting time required
- Otherwise, synchronized time become slower wrt external clock
- Slaves still remain synchronized within a specified bound
- Very difficult to compensate these in practical systems

## Example 2

Assume that the drift rate between any two clocks is restricted to  $\rho = 5 \times 10^{-6}$ . Suppose we want to implement a synchronized set of six distributed clocks using the central synchronization scheme so that the maximum drift between any two clocks is restricted to  $\epsilon = 1$  msec at any time, determine the period with which the clocks need to be resynchronized.

- Solution. The maximum drift rate between any two arbitrary clocks when the clocks are synchronized using a central time server with a resynchronization interval of  $\Delta T$  is given by  $2\rho\Delta T < \epsilon$ . Therefore, the required resynchronization interval  $\Delta T$  can be expressed as:


$$\Delta T < \frac{1 \times 10^{-3}}{5 \times 10^{-6} \times 2} \text{ sec} = \frac{10^{-3}}{10^{-5}} \text{ sec} = \frac{1}{10^{-2}} \text{ sec} = 100 \text{ sec.}$$

Therefore, resynchronization period must be less than 100 sec.

# Distributed Clock Synchronization


- Centralized clock synchronization susceptible to single point failure
- Any failure of master clock causes breakdown of synchronization
- Distributed clock synchronization - no master clock wrt all slaves
- All clocks periodically exchange clock readings among themselves
- Based on received time, clocks compute and set synchronized time
- Possible that some clocks are bad or become bad during operation
- Bad clocks exhibit larger drifts than specified tolerance
- Bad clocks may even stop keeping time

# Distributed Clock Synchronization



- Bad clocks can be identified and taken care of during synchronization
- By rejecting time values of any clock larger than specified bound
- Usage of Byzantine clocks
- A Byzantine clock is a two face clock
- Transmits different values to different clocks at same time

# Byzantine Clock



- $C_1$  is a Byzantine clock
  - sending time value  $t+e$  to clock  $C_5$
  - $t-e$  to clock  $C_3$
  - at same time instant
- If  $< \frac{1}{3}$  clocks bad (Byzantine), good clocks approx. synchronized
- Synchronization scheme of clocks
  - Let there be  $n$  clocks in a system
  - Each clock periodically broadcasts time value at end of certain interval
  - Assume clocks to be synchronized within  $\epsilon$  time units of each other
  - If a clock receives broadcast time differs from its own time by  $> \epsilon$
  - then sending clock must be bad, safely ignores received time value
  - Each clock averages all good time values received
  - sets its time with average

# Pseudo Code for Distributed Clock Synchronization

- Each clock  $C_i$  carries out following operations

Procedure distributed clock synchronization:


```
good-clocks=n;  
for(j=1;j < n;j++){  
    if ( $|(c_i - c_j)| > \epsilon$ ) good-clocks--; // Bad clock  
    else total-time= total-time +  $c_j$ ;  
     $c_i$ =total-time/good-clocks; // set own time equal to the computed time  
}
```

- Each clock carries same set of steps
- If all  $n$  clocks carry these steps, then atmost  $m$  are bad
- $n > 3m$
- Good clocks will be synchronized within  $3\epsilon m/n$  bound



# Theorem 1

In a distributed system with  $n$  clocks, a single Byzantine clock can make two arbitrary clocks in a system to differ by  $3\epsilon/n$  in time value, where  $\epsilon$  represents the maximum permissible drift between two clocks.



- Proof. Consider three clocks C1, C2 and C3 in a distributed system. C1 is a Byzantine clock. C2 and C3 are good clocks - differ by  $\neq \epsilon$ . C3 being Byzantine shows two different values to C1 and C2. Effect of Byzantine clock in total time calculation makes two good clocks differ by at most  $3\epsilon$ . Effect of a single Byzantine clock make two arbitrary clocks differ by  $3\epsilon/n$ . For  $m$  Byzantine clocks two good clocks differ by at most  $3\epsilon m$  in average. Individual clocks synchronized within  $3\epsilon m/n$ .

# Time required for clocks drift to $\epsilon$

- Let time required for two clocks to drift from  $3\epsilon m/n$  to  $\epsilon$  be  $\Delta T$ .  
or,  $2\Delta T\rho \leq (n\epsilon - 3\epsilon m)/n$   
or,  $\Delta T \leq (n\epsilon - 3\epsilon m)/(n \times 2\rho)$   
 $\Delta T$  is time required for two good clocks to drift from  $3\epsilon m/n$  to  $\epsilon$   
 $\Delta T \leq [(3m+1)\epsilon - 3\epsilon m]/(n \times 2\rho)$   
 $\Delta T \leq \epsilon/2n\rho$

## Example 3

Let a distributed real-system have 10 clocks, and it is required to restrict their maximum drift to  $\epsilon = 1$  msec. Let the maximum drift of the clocks per unit time ( $\rho$ ) be  $5 \times 10^{-6}$ . Determine the required synchronization interval.

- Solution.  $\Delta T = \frac{10^{-3}}{2 \times 10 \times 5 \times 10^{-6}}$

$$\Delta T = 10 \text{ sec}$$

Required synchronization interval is 10 sec.

# Thank you